

UNIVERSITY OF MANCHESTER

Mandelbrot on FPGA

by

Obaya Dralega

A thesis submitted in partial fulfillment for the
degree of Bachelor of Science

in the
School of Computer Science

May 2016

UNIVERSITY OF MANCHESTER

Abstract

School of Computer Science

Bachelor of Science

by Obaya Dralega

The Mandelbrot Set is one of the most popular fractal images today. The term fractal is applied to mathematically defined shapes or objects that display self-similarity at various scales. Although mathematically defined, fractal shapes are often seen in the natural world, from trees to clouds to coastlines.

FPGAs (Field Programmable Gate Arrays) are computing devices composed of configurable logic blocks connected by a programmable Interconnection Fabric. An FPGA is able to be reprogrammed to implement any digital circuit with its capacity. This means that an FPGA is able to implement a Hardware overlay, which is an FPGA implementation of an existing processor architecture.

This project report follows the progress in the research, development and implementation of a system for the computation of the image of the Mandelbrot set on an FPGA Overlay. Due to several unforeseeable delays to my progress, as well as a steep learning curve, I was unable to implement my final design on an FPGA, although I proved my design functionally correct through a software implementation and simulation. Screen-shots of my software results are presented along with a summary of the implementation details of my design.

Acknowledgements

I would like to thank my supervisor, Dr. Dirk Koch for his endless enthusiasm and invaluable advice.

Contents

Abstract	i
Acknowledgements	ii
List of Figures	v
1 Introduction	1
1.1 Mandelbrot Set	1
1.1.1 Mathematical Description of the Mandelbrot Set	2
1.1.2 Computation of the Mandelbrot Set	3
1.1.3 Zooming into the Mandelbrot Set	4
1.1.4 Computational considerations when Computing the Mandelbrot Set	5
1.1.4.1 Finite Resolution	5
1.1.4.2 Accuracy versus Time	5
1.2 FPGA	6
1.2.1 FPGA Architecture	6
1.2.2 Why use FPGAs ?	7
1.2.2.1 Mandelbrot and FPGAs	8
1.2.3 Hardware Platform	8
1.3 FPGA Programming	9
1.3.1 Hardware Description Languages	9
1.3.2 Schematic Design	9
1.3.3 Hardware Overlays	9
1.4 MIPS Architecture	11
1.4.1 MIPS Functional Blocks	11
1.4.2 MIPS Instruction Set	12
1.5 Number Representations	14
1.5.1 Floating Point	14
1.5.2 Fixed Point	14
2 Design and Implementation	16
2.1 System Implementation Overview	16
2.2 Mandelbrot C code	17
2.3 Complex Number C code	17
2.4 Fixed Point Code	19

2.4.1	Comparison of Fixed and Floating Point Representations for Mandelbrot	19
2.4.2	Scale Factor	20
2.4.2.1	Effect of varying the Scale Factor	20
2.4.3	Conversion	22
2.5	Cross Compilation	23
2.5.1	Building the Cross Compiler Toolchain	23
2.6	Effects of Number of Iterations on Image Quality	24
3	Evaluation	26
3.1	Code Evaluation	26
3.2	Simulation	27
4	Reflections and Conclusion	29
4.1	Change of Objective	29
4.2	Challenges	29
4.3	Conclusion	30
5	Further Research	31
5.1	The Mandelbulb	31
5.1.1	Changes to be Made	32
A	Run Away to Infinity Criterion	33
	Bibliography	35

List of Figures

1.1	The full Mandelbrot set image.	1
1.2	A complex number can be visually represented as a pair of numbers forming a vector on a diagram called an Argand diagram, representing the complex plane.	2
1.3	A visual representation of the magnitude of a complex number.	3
1.4	A zoom into a section of the Mandelbrot set.	4
1.5	A diagram showing how the centre points of pixels are chosen from a finite pixel array to represent complex numbers [3].	5
1.6	Generic FPGA Logic Element [5].	6
1.7	Generic Island Style FPGA Architecture [5].	7
1.8	The Xilinx Nexys 4 DDR FPGA Board [11].	8
1.9	Overview of the MIPS functional components [18].	12
1.10	MIPS instruction types [18]	13
1.11	How floating point numbers are represented at the bit level [22].	14
1.12	How fixed point numbers are represented at the bit level [23]	15
2.1	A diagram giving a high-level design overview of the system.	16
2.2	Mandelbrot render with scale factor 4.	21
2.3	Mandelbrot render with scale factor 8.	21
2.4	Mandelbrot render with scale factor 24.	21
2.5	Macros for handling fixed point numbers.	22
2.6	Unmodified main loop body of Mandelbrot computation using floating point numbers.	22
2.7	Modified main loop body of Mandelbrot computation using fixed point numbers.	22
2.8	A magnified section of the Mandelbrot set image computed using an increasing number of maximum of iterations.	25
3.1	A screen shot of the some of the cross-compiled assembly code in the the Instruction BRAM of the MIPS overlay.	27
3.2	A screen shot showing the button which launches the simulator after checking the Behavioural Syntax.	27
3.3	A screen shot of the simulation window.	28
3.4	A close up screen shot of the most relevant signals in the simulation. . . .	28
5.1	A render of the Mandelbulb [31].	31

Dedicated to my parents, Audrey and Anguyo Dralega.

Chapter 1

Introduction

This chapter will introduce the key concepts that are needed to understand the technical details of the project and describe the project in a broad context.

1.1 Mandelbrot Set

The Mandelbrot Set is a fractal named after the mathematician Benoit Mandelbrot. Fractals are objects that display self-similarity at various scales [1]. Magnifying a fractal reveals small-scale details similar to the large-scale characteristics. Although the Mandelbrot set is self-similar at magnified scales, the small scale details are not identical to the whole. In fact, the Mandelbrot set is infinitely complex, yet the process of generating it is based on a simple equation involving complex numbers.




FIGURE 1.1: The full Mandelbrot set image.

From a philosophical perspective, the Mandelbrot set challenges familiar notions of simplicity and complexity: how could such a simple formula, involving only multiplication and addition, produce a shape of great organic beauty and infinite subtle variation? It defies conventional intuition and has captured the imagination of Computer Scientists and Mathematicians for decades.

1.1.1 Mathematical Description of the Mandelbrot Set

The Mandelbrot set is mathematically defined as the set of complex numbers for which the sequence $z_{n+1} = z_n^2 + c$ does not approach infinity when applied iteratively [2].

$$z_{n+1} = z_n^2 + c$$

$$c \in M \iff \lim_{n \rightarrow \infty} |z_{n+1}| \leq 2$$

Complex numbers consist of a real number added to an imaginary number. While real numbers can be represented on one dimensional number line, complex numbers have two parts, so they are represented on a 2-dimensional plane called the complex number plane or *Argand diagram* as can be seen below.

$$\text{ComplexNumber} : z = x + yi$$




FIGURE 1.2: A complex number can be visually represented as a pair of numbers forming a vector on a diagram called an Argand diagram, representing the complex plane.

1.1.2 Computation of the Mandelbrot Set

To compute the Mandelbrot set we need to calculate which pixels of the display belong to the set, we do this by applying the formula, $z_{n+1} = z_n^2 + c$ iteratively to each pixel. We start with $z_0 = 0$ and c takes its value from the co-ordinates of the pixel in the complex plane.

When graphing the Mandelbrot set, each pixel in the screen's display is taken for a co-ordinate of this complex plane. We need to map the pixel's screen coordinate to coordinates in the complex plane. The y-axis co-ordinate representing the imaginary component and the x-axis the real component of the complex number. Which complex coordinates we need in particular depends on which part of the set we are looking at.

When the Mandelbrot set is computed, we are not too interested in the actual value of z , we are interested in the magnitude of z , which is its distance from the origin. This can be seen in Figure 1.3. For a complex number $z = x + yi$, the magnitude r is calculated by taking the square root of the sum of the real and imaginary components squared.

$$z = x + iy$$

$$|z| = r$$

$$r = \sqrt{x^2 + y^2}$$




FIGURE 1.3: A visual representation of the magnitude of a complex number.

If the magnitude of z , for a given pixel on screen, does not exceed 2 after a certain number of iterations, then it is likely to remain bounded. The pixel belongs to the Mandelbrot set and is coloured black. If the magnitude of z , for a pixel on screen, exceeds 2, then it will eventually tend to infinity (Due to the Run Away to Infinity Criterion. The theorem's proof can be seen in Appendix A) and the pixel is coloured based on how many iterations it took for the formula to become unbounded. It should be noted that each pixel is computed independently of all other pixels.

The behaviour of the magnitude is analogous to squaring a number. Iteratively squaring a number greater than 1 will cause it to grow forever, however, an iteratively squared number less than 1 will never grow large, it will get smaller.

1.1.3 Zooming into the Mandelbrot Set

The scale at which we see the Mandelbrot set image is dependent on the size of the area of the image we are looking at. The zoom level of the image can be adjusted by adjusting the mapping of the complex plane area to screen area. This ratio is all that is needed to be changed. The screen area is fixed (as it is a discretely finite array of pixels), but we can zoom in by specifying that a smaller area of the infinite complex plane should be mapped to it. Figure 1.4 shows a zoom into an area of the Mandelbrot set, revealing the characteristic Mandelbrot swirls as well as a smaller copy of the whole image.




FIGURE 1.4: A zoom into a section of the Mandelbrot set.

1.1.4 Computational considerations when Computing the Mandelbrot Set

1.1.4.1 Finite Resolution

Practically, when computing the Mandelbrot set, we only consider a portion of the complex plane, the window in which the picture is rendered. A computer monitor's display is actually a finite and discrete array of pixels. The complex plane, however, is continuous and contains an infinite number of points. It is not possible to resolve an image at a level smaller than a pixel. Due to this, one point in the pixel is taken to represent a coordinate in the complex plane, usually the centre of the pixel [3]. Figure 1.5, visually demonstrates this.




FIGURE 1.5: A diagram showing how the centre points of pixels are chosen from a finite pixel array to represent complex numbers [3].

1.1.4.2 Accuracy versus Time

If, for a given pixel, all of z_0, \dots, z_n lie within a distance of 2 from the origin (have a magnitude less than 2) for relatively large values of n , we may conclude that the sequence does not run to infinity. However, this may not be necessarily true. If, for a given pixel, all of z_0 to z_{1000} lie within a distance of 2 from the origin, but z_{1001} lies outside we may incorrectly classify a point as belonging to the Mandelbrot set, when it does not. But how many iterations are enough? We must make a choice and select some maximum number of iterations we are willing to try.

As the number of iterations increases, the number of wrong classifications decreases and the image will have finer detail. Unfortunately, increasing the number of iterations also increases the computer time needed to generate the picture.

1.2 FPGA

This section will give a high-level overview of FPGAs, their architecture, how they function and the hardware device that my implementation is designed for.

FPGAs (Field Programmable Gate Arrays) are programmable digital chips [4]. They can be thought of as *programmable hardware*, as they are able to implement an arbitrary number of digital circuits within the device's capacity. They are flexible in use and are designed to be configured by the user after their manufacture.

1.2.1 FPGA Architecture

FPGAs are composed of large numbers of small logic blocks called *Logic Cells*, which perform the general logic computations. Logic Cells are typically composed of a 6-input Lookup table which can act as a small RAM, as can be seen in Figure 1.6. When the FPGA is configured (the initial programming of the FPGA to set the chosen behaviour of the logic cells), each logic cell is loaded with the 2^6 bits holding the Boolean truth table of the particular logic function we want to implement [5].

Logic cells are grouped to form *Logic Clusters* which are able to compute logic functions of greater complexity. The Logic clusters are then arranged regularly in the *Interconnect Fabric* in an *Island-Style architecture* [6], as can be seen in Figure 1.7. This enables the connection of and communication between multiple Logic Clusters.




FIGURE 1.6: Generic FPGA Logic Element [5].

Along with Logic Clusters, the Island-Style architecture also includes blocks of logic which implement specific functions efficiently. Block RAM (BRAM) is used for larger and faster storage while Digital Signal Processing (DSP) Blocks are used for quicker arithmetic functions such as multiplication. In addition to these, FPGAs also contain other building blocks such as flip flops and I/O blocks.




FIGURE 1.7: Generic Island Style FPGA Architecture [5].

1.2.2 Why use FPGAs ?

The key to the usefulness of an FPGA lies in its ability to be customised. With FPGAs, the machine can be tailored to needs of the application, giving significant speed ups over software implemented on CPU or GPUs [7].

In terms of performance, FPGAs excel on stream processing problems as data movement can be synchronised to minimise data movement on the chip. Furthermore algorithms can be implemented directly after each other in a chain, eliminating the need for off-chip communication which provides further performance benefits [8].

Practically, FPGA are used in various domains, from network equipment, to avionics, to medical devices and data processing [9].

1.2.2.1 Mandelbrot and FPGAs

We have seen that when computing the Mandelbrot set each pixel is independently iterated on with no data dependencies between pixels. This shows us that the Mandelbrot Set is an inherently parallelisable and a purely compute-bound problem [10]. It is ideally suited for computation on FPGAs as an FPGA can take advantage of horizontal scaling where multiple threads running side by side implement the same algorithm but with different data [8].

1.2.3 Hardware Platform

For this project I used the Nexys 4 DDR Artix-7 from XILINX [11]. This board, pictured below, is equipped with 240 DSP slices, 4,860 Kbits of BRAM, 15,850 slices each with four 6-input LUTs and eight flip-flops.




FIGURE 1.8: The Xilinx Nexys 4 DDR FPGA Board [11].

1.3 FPGA Programming

This section will detail the various ways in which an FPGA can be programmed. There exist several methods to define the behaviour of an FPGA, each with its respective advantages and disadvantages.

1.3.1 Hardware Description Languages

Hardware Description Languages (HDLs), such as Verilog and VHDL [12], are specialised text-based computer languages used to describe the structure and behaviour of digital circuits. They are precise and permit the description of low level details of the hardware relatively directly. HDLs were first created to implement a register-transfer level abstraction [13] (a model of the data flow and timing of a circuit). Viewed in this way, HDLs can be compared to the level of abstraction of assembly languages.

Hardware Description Languages are good for smaller circuits, but do not scale well for larger designs as any high-level components must be detailed at a low level [5, Section 2.2.1]. This can make the development process unnecessarily complex.

1.3.2 Schematic Design

Schematic Design makes use of a visualisation of the system, it uses abstract graphical symbols to represent the system components [14]. While linking boxes with wires on a screen appears to be much simpler than specifying each component of a system in text, in practice, it can become more complex. Available schematic editors often have poor user interfaces. Connecting large numbers of wires and buses on a 2-dimensional page can quickly become confusing and incomprehensible for larger designs. Furthermore, schematic design does not permit the programmer to express loop constructs, which are supported in HDLs.

1.3.3 Hardware Overlays

Hardware Overlays are distinct from HDLs and Schematic Design in that they provide a layer of abstraction over the FPGA hardware. As mentioned earlier, FPGAs are flexible and can be programmed to implement any digital circuit within their capacity. In theory this means that an FPGA can be used to implement another programmable architecture on top of the FPGA fabric, such as an x86 CPU for example. The implemented

architecture can then be programmed using software alone. Any software that runs on an actual x86 processor can be directly migrated to the x86 overlay [15], [16].

A major advantage to the hardware overlay approach is that implementing a new or updated system does not require the use of CAD (Computer Aided Design) tools. This eliminates the significant amount of time spent on compilation and Place and Route (The placement and interconnection of logic elements on the FPGA). Furthermore, a system that has been implemented using an overlay is portable across separate implementations of the same overlay on different devices [5].

However, there are costs associated with the overlay. Performance suffers, both when compared to the system implemented directly on the underlying hardware and when compared to the software running on on a Hard CPU [17].

1.4 MIPS Architecture

This section will give context about the design of the underlying MIPS system that my design is to be built on including details the functional design as well as the instruction set. For this project my system was built upon a base of the MIPS processor overlay designed by Anders Hauk Fritzell in his Master's thesis *A system for Fast Dynamic Partial Reconfiguration using GoAhead* [18].

MIPS is an acronym for *Microprocessor without Interlocked Pipeline Stages*. It is a popular choice in the embedded systems industry, included in devices such as routers and video game consoles [19]. MIPS is a Reduced Instruction Set Computer (RISC) architecture. This means that its Instruction Set Architecture (ISA) is made up of small, fixed-length, heavily optimised instructions that are executable in a single cycle. These instructions are straightforward to decode due to their fixed length format. Another characteristic trait of MIPS is that memory accesses are exclusively limited to the load and store instructions and all operations are done within the registers of the microprocessor [20].

1.4.1 MIPS Functional Blocks

The MIPS processor is composed of several functional blocks as pictured below in Figure 1.9 [18]:

- **Instruction Decoder:** Takes instructions and decodes them.
- **Program Counter (PC):** Holds the address of the next instruction to be executed.
- **Arithmetic Logic Unit (ALU):** A main component that performs the arithmetic and logical operations on data.
- **Registers:** 31 general purpose registers , r0 holds a constant zero.
- **Memory:** Holds the data values to be used. Only accessible through load and store instructions.

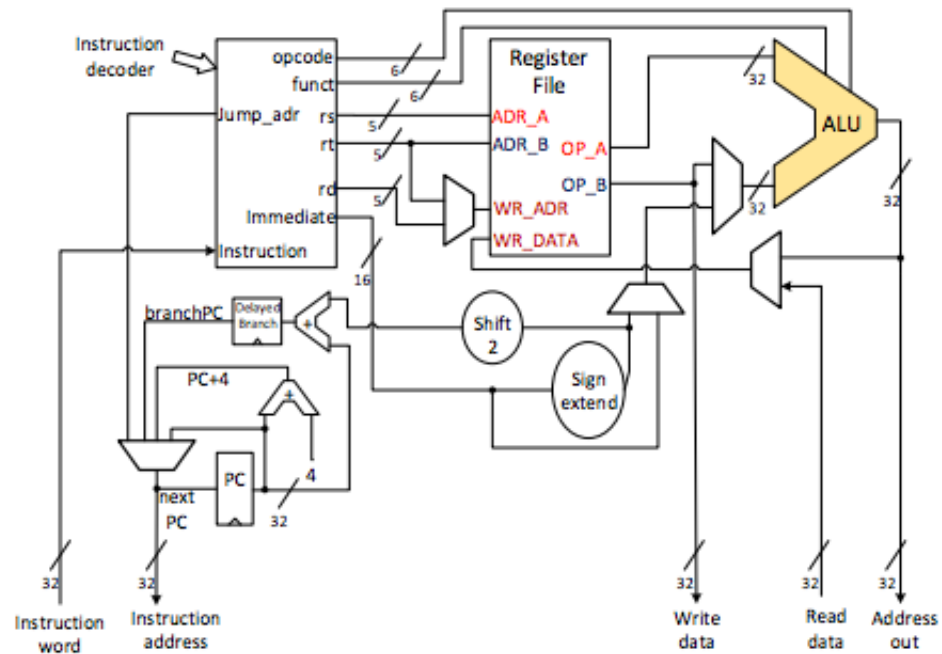


FIGURE 1.9: Overview of the MIPS functional components [18].

1.4.2 MIPS Instruction Set

The MIPS architecture has three 32-bit instruction types [18]:

- **R-type** (Register Format): Register to Register. Operands are taken from the 5-bit register specifications *rs* and *rt*. *rd* is a destination register for the result. The type of operation is determined by *funct* (6-bits) as many R-Type instructions share an *opcode*.
- **I-type** (Immediate Format): The operation is performed with a constant 16-bit value *immediate*. The operand is taken from *rs* and result is put in *rt*. This instruction type is also used for branching. The immediate is added to the PC to perform a PC-relative branch.
- **J-type** (Jump Format): Used to provide a new address to the program counter, moving execution to a new code block.

Type	Bits					
	31-26	25-21	20-16	15-11	10-6	5-0
R	opcode	rs	rt	rd	shamt	funct
I	opcode	rs	rt	immediate		
J	opcode	address				

FIGURE 1.10: MIPS instruction types [18]

In the above table:

- **opcode** (Operation Code): How an instruction should be decoded.
- **rs,rt**: Addresses of source register operands.
- **rd**: Destination register operand.
- **shamt** (Shift Amount): Used in shift instructions.
- **funct**: Select the variant of the operation in the opcode field.
- **Immediate**: 16-bit constant used for constant value operations.
- **Address**: used in jump instructions to move execution to another part of code.

The MIPS instructions have been designed in this way purposefully. The ISA exhibits some key underlying principles in the design of hardware.

1. Regularity helps with Simplicity. For example, the two arithmetic instructions *add* and *subtract* both have exactly three operands. This simplifies the design of the ALU.
2. Smaller is faster, for this reason MIPS only has 32 registers. Having greater numbers of registers would lead to an increase in clock cycle time.
3. Make the common case fast. The common case is the case used most frequently. Increasing the speed of this single option can lead to even greater increases in overall speed. MIPS addresses this maxim by incorporating constants as part of the arithmetic instructions and loading small constants into the upper 16-bits of a register.

1.5 Number Representations

There are infinitely many real numbers, however a computer is a finite state machine and thus, can only store a finite amount of information. As numbers cannot be stored with infinite precision, representing real numbers in a computer always involves an approximation [21]. This section will explain the 2 main methods of real number representation on a computer, namely fixed point and floating point.

1.5.1 Floating Point

Floating point numbers are similar in concept to scientific notation. Floating point gains its name for the fact that the number's radix point can *float*. This means that the radix point can be placed anywhere relative to the a numbers significant digits.




FIGURE 1.11: How floating point numbers are represented at the bit level [22].

Floating point numbers are composed of two main parts:

- **The Mantissa:** (or Fraction) Contains the number's digits.
- **The Exponent:** Defines where the decimal or binary point is placed relative to the beginning of the Mantissa. Negative exponents represent numbers less than 1.

1.5.2 Fixed Point

A fixed point number is actually an integer that is scaled by a chosen factor. The integer values involved in some arithmetic do not represent their normal integer values but instead represent their values with a radix point at some *fixed* location. The scaling factor essentially determines the position of the radix point and is fixed for the duration of the entire computation.

This concept is most easily understood by examples, in this case with decimal numbers.

If I want to represent the real, fractional value 1.67, I can instead use the value 1670 with an applied scaling factor of 1/1000. Conversely, the number 252,000 can be represented as 252 with a scaling factor of 1000.




FIGURE 1.12: How fixed point numbers are represented at the bit level [23]

An example of fixed point arithmetic is the multiplication of $0.0456 * 2.789 = x$. First we multiply both numbers by 10000.

$$456 * 27890 = y$$

Now both numbers are integers.

$$y = 12717840$$

We have an value for y, but we originally wanted a value for x.

$$x = y/10000^2$$

$$x = 0.1271784$$

This is the correct answer computed using only integer arithmetic.

Chapter 2

Design and Implementation

This section will detail the design approach to the project, going into detail on how the Mandelbrot Set is computed in code. First, an overview of the implemented system will be given, then the step-by-step details of the implementation process.

2.1 System Implementation Overview

The Mandelbrot set image computation is handled by C code which takes pixel screen coordinates from an address, computes the colour of the pixel, then returns this colour to an output address. The C code is then cross compiled into bare-metal instructions. The bare-metal instructions are then modified to the correct format and placed in the Instruction ROM of the MIPS overlay. The code is then run on the MIPS CPU overlay on the FPGA, showing the Mandelbrot Set on screen. This is seen in Figure 2.1.

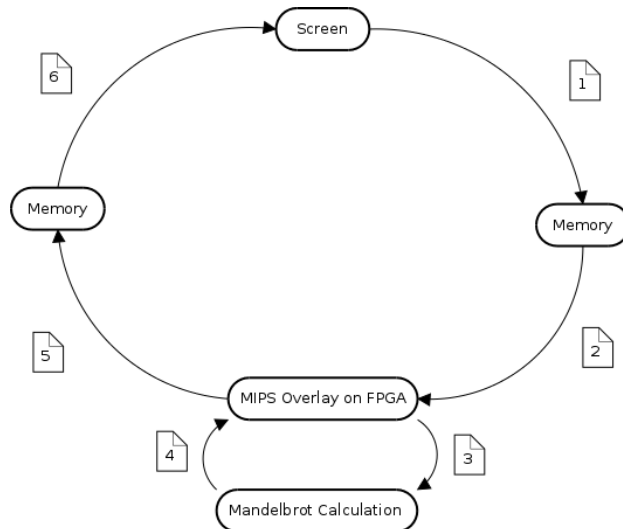


FIGURE 2.1: A diagram giving a high-level design overview of the system.

2.2 Mandelbrot C code

The C code can be seen as the starting point of the implementation. The code first takes pixel screen coordinates from an input address, the inputted pixel coordinates are then passed to a Mandelbrot computation function. In this function the pixel coordinates on the screen are mapped to coordinates in the complex plane. The pixel is then iteratively checked to see if it belongs the Mandelbrot set in a loop. If the pixel does belong to the set, then the colour black is returned. If not, a colour based on the number of iterations it took to break out of the loop is returned.

The difficulty in writing the code came from two areas which are discussed below. The first difficulty came from the fact that complex numbers are not directly represented in C and the code would need to be implemented using integers. The second difficulty came from the fact that floating point arithmetic was unsuitable for my implementation as the MIPS overlay did not have a Floating Point Unit (FPU), so the code needed to be modified to use fixed point arithmetic.

2.3 Complex Number C code

This section will show how complex number arithmetic can be computed using a series of integer calculations [24].

In the Mandelbrot algorithm, we have two complex numbers we make use of : Z_n and c . As mentioned earlier in the Section 1.1.1 of this report, complex numbers are made up of a real and imaginary part.

As complex numbers have no native type in the C language, we need to find away to represent them in another way. This can be done using elementary algebra and integers to represent the imaginary and real components of the complex number.

$$z_n = x + iy \tag{2.1}$$

$$c = a + ib \tag{2.2}$$

From the Mandelbrot algorithm we want to compute z_{n+1} which is the next iteration :

$$z_{n+1} = z_n^2 + c$$

Below is the expansion of z_n^2 :

$$\begin{aligned}
 z_n^2 &= (x + iy)^2 \text{ [substituted from (2.1) above]} \\
 &= x^2 + i2xy + i^2y^2 \text{ [expansion of brackets]} \\
 &= x^2 + i2xy - y^2 \text{ [as } i^2 = -1 \text{ by definition of imaginary numbers]} \\
 &= (x^2 - y^2) + i2xy \text{ [finally, combine the real and imaginary parts]}
 \end{aligned}$$

The end result is also a complex number as it is the addition of real and imaginary components. To complete the expression, we now add the complex number c to get $z_{n+1} = z_n^2 + c$.

$$\begin{aligned}
 z_{n+1} &= (x^2 - y^2) + i2xy + c \\
 &= (x^2 - y^2) + i2xy + a + ib \text{ [substituted from (2.2) above]} \\
 &= (x^2 - y^2 + a) + i(2xy + b) \text{ [combine real and imaginary parts]}
 \end{aligned}$$

Again, we can see the end result is also a complex number, but it has been split into parts and so can be calculated as a series of integer calculations.

2.4 Fixed Point Code

This section will contrast fixed and floating point number representations, showing that for that for this implementation, fixed point is the better choice. Following this, the implementation details of fixed point will be presented.

2.4.1 Comparison of Fixed and Floating Point Representations for Mandelbrot

An advantage of floating point is that the floating point exponentiation allows us to represent numbers with large dynamic ranges (great differences in magnitude). This is important in cases where the data set is large or the range of the data is unpredictable.

A further advantage of floating point is that most modern programming languages include a floating point data type, for example the *double* type in C. Generally speaking, this makes it simpler to implement a complex algorithm using the floating point type. Fixed point has no native data type and the code must be modified to accommodate them.

However there are some drawbacks to floating point. If you know ahead of time that the numbers you will be using for a given program fall within a small range, then floating point may not be useful due to the way floating point numbers are represented. If the numbers are so small that they have a maximum exponent of 1, then the 6-bits reserved for the exponent are unused and inaccessible. This is the case in the Mandelbrot computation where all of the numbers used fall with a small range of $[-2, 2]$.

Fixed point does not have the disadvantage of wasted bits as the place of the radix point can be wholly determined by the programmer.

An advantage of fixed point is that implementing an FPGA Overlay which only has support for fixed point integer arithmetic makes the implemented design substantially smaller. This has the follow on effect of making the design less expensive in terms of resources consumed on the FPGA board and also faster.

Another advantage of fixed point is its flexibility of use. The precision, range and data size can all be chosen by the programmer. This makes fixed point suited to implementations that require specific data formats. The freedom to choose the precision is crucial for zooming into the Mandelbrot set image, where at deep zoom levels several hundred bits of accuracy are needed.

The major drawback to fixed point is the limited range of representable numbers. The maximum and minimum fixed point numbers are determined by the maximum and minimum numbers that can be stored by underlying data type multiplied by the scaling factor.

It is a common misconception that floating point number representations are generally preferable for the computation of the Mandelbrot set. While floating point representations have some advantages, when computing the Mandelbrot set, fixed point is ultimately the winner. As well as this, the MIPS overlay I used in my implementation did not have a Floating Point Unit (FPU) and so could not support floating point arithmetic.

2.4.2 Scale Factor

In order to gain the extra accuracy gained by use of fixed point number representations, modifications need to be made to the code.

The scale factor determines the position of the radix point, so when the scale factor is chosen, we are selecting the range and precision of the arithmetic that will be performed in our program. As we are performing the arithmetic on a binary computer, we choose a scale factor that is a power of 2 as binary computers are able to quickly multiply and divide by powers of 2 using logical bit shifting.

For this reason I use the scale of $1/2^{\text{Shift_Factor}}$. I set a constant variable, *Shift_Factor*, and use it for shifting.

This program uses 32-bit integers, I have chosen a scale factor of 24, meaning that, in effect, there are 24 decimal places after the radix point and all values will be multiplied by 2^{24} before their use in any arithmetic.

2.4.2.1 Effect of varying the Scale Factor

A scale factor of 24 has been chosen after experimentation with different scale factors. As can be seen below in Figures 2.2-4, as the scale factor increases the image quality of the Mandelbrot Set increases. This is to be expected, as with very few decimal places available, the coordinates passed to the Mandelbrot computation function are less accurate.

The first image, with a scale factor of 4, is of poor quality and is pixelated, although it still retains the characteristic Mandelbrot set shape. When the scale factor is increased to 8, there is a significant jump in image quality. When further increased, the improvement

in image quality is much less noticeable at this scale, but upon zooming in finer detail than the previous images is visible.

All these images were computed by the Mandelbrot test program using a constant resolution of 2560 x 1440.




FIGURE 2.2: Mandelbrot render with scale factor 4.




FIGURE 2.3: Mandelbrot render with scale factor 8.




FIGURE 2.4: Mandelbrot render with scale factor 24.

2.4.3 Conversion

This subsection shows how different functions are used to convert between fixed point integer representations and normal integers.

```
6 #define multiplyFixed(a,b) (((long long)a)*(b))>>SHIFT_AMOUNT)
7 #define FixedPointConvert(a) ((long)(((a)*(1<<SHIFT_AMOUNT))))
8 #define integer(a) (((a)+(1<<(SHIFT_AMOUNT-1)))>>SHIFT_AMOUNT)
```

FIGURE 2.5: Macros for handling fixed point numbers.

Above we can see the macro functions that are used for fixed point arithmetic.

MultiplyFixed(): takes two large integers as input and multiplies them together. The product is then rescaled by right-shifting it by the shift amount.

FixedPointConvert(): takes a Fixed Point integer value as input, converts it to a long integer and multiplies it by 1 shifted by the scale factor to produce the output. This is equivalent to multiplying the number by 2^{Shift_Amount} .

integer(): takes a long integer value and converts it into a small fixed point number by adding it to $2^{Shift_Amount-1}$, then rescaling the sum by right-shifting by the shift amount.

```
80 while((x*x + y*y < 4) && iteration < MAX_ITERATION)
81 {
82     xtemp = x*x - y*y + x0;
83     y = (2*(x*y)) + y0;
84     x = xtemp;
85     iteration++;
86 }//while
```

FIGURE 2.6: Unmodified main loop body of Mandelbrot computation using floating point numbers.

```
while ((multiplyFixed(xn,xn)+multiplyFixed(y0,y0))<FixedPointConvert(4) && ++i<MAX_ITERATIONS)
{
    xn=multiplyFixed((x0+y0),(x0-y0)) +p;
    y0=multiplyFixed(FixedPointConvert(2),multiplyFixed(x0,y0)) +q;
    x0=xn;
}//while
```

FIGURE 2.7: Modified main loop body of Mandelbrot computation using fixed point numbers.

In Figures 2.6 and 2.7 we see the comparison of the original floating point implementation and the fixed point implementation of the main loop of the Mandelbrot computation. The different implementations are almost identical in structure, but in the fixed point implementation most arithmetic performed on the values is handled by the macro functions defined in Figure 2.5.

As can be seen in Figure 2.7, the defined macros can then be simply substituted into the code to make it fixed point. An advantage of this approach is that the same fixed point macros can be used to convert other programs to fixed point with little additional work. As previously stated, the MIPS overlay used in this design does not include a Floating Point Unit, so this feature is necessary when expanding the design to compute an image other than the Mandelbrot set.

2.5 Cross Compilation

This section will cover one of the central elements of this systems design, the method by which the C code is able to run on the overlay. It will give a brief overview of compilers, then provide some detail on the implementation of the cross compiler used in this project.

A compiler is a program that generates executable code from source code. Typically, what is colloquially referred to as a *compiler* is actually a *native compiler*. A native compiler runs on a specific type of computer and generates code to also run on that same specific type of computer.

A *cross compiler* is a compiler that runs on one platform (called the host) but generates executable code for another, different platform (called the target) [25]. The two platforms may have different operating systems and even different CPUs. An easily understood example would be a Microsoft Windows compiler generating code for Apple iOS device.

Cross compilation has several applicable uses, such as embedded systems design. The embedded device is unlikely to have the computing resources to run a development environment or a compiler. Another example is Bootstrapping, a process in which a compiler is written in the source programming language which it intends to compile.

2.5.1 Building the Cross Compiler Toolchain

For this project I needed a compiler which would take C code and output a bare metal executable. I had a choice of two cross compilers, the *mips-linux-gnu* and the *mipsel-elf*. Both of these compilers met the above requirement, the decision came down to using both compilers to compile a short test program, then using *objdump* (A program that displays information about object files) to view the assembly code and comparing which came closer to my desired output. I chose the 'mips-linux-gnu' cross compiler as the compiler's target platform is bare-metal MIPS, its host platform is a linux machine.

2.6 Effects of Number of Iterations on Image Quality

This section shows some of the experimentation that was carried out in order to determine the optimum number of iterations to set for the final implementation.

As mentioned in Section 1.1.4.2, there is a trade off between the accuracy of the image and the time spent computing the image. A choice needs to be made on the number of iterations before it is decided that a pixel does or does not belong to the Mandelbrot set. Before running the simulation, I experimented with what the minimum number of iterations was to see acceptable levels of detail.

The main factor in this case is image quality, as there is nothing gained by fast computation if the image is of poor quality.

Below we can see the Mandelbrot images that is computed using different maximum numbers of iterations. The colour yellow indicates the points that the Mandelbrot computation has decided belong to the Mandelbrot set.


In Figure 2.8a, when the maximum number of iterations is limited to 128, the image quality is extremely poor and "blob-like". There is almost no fine detail visible and bands of colour can be clearly seen. The large yellow region shows that at such a low number of iterations there have been many wrong classifications. However, even at this low number the overall outline is still correct.

When the maximum number of iterations doubles to 256 in Figure 2.8b, some fine detail is now visible but the bands of colour can still be seen and the centre regions still appear blob-like.


Figure 2.8c makes a jump in the maximum number of iterations to 896. It is very detailed and clear. No bands of colour can be seen and the image resembles particles rather than blobs.

Past 896, an increase in the maximum number of iterations has little effect. In Figure 2.8d the maximum number of iterations is set at 3768. This is a significant jump up from 896 but the images are almost identical.


For my implementation I chose a final maximum number of iterations of 896 as beyond this point there is no image improvement to be gained but the time to compute the image increases.




(a) Max Iterations: 128



(b) Max Iterations: 256



(c) Max Iterations: 896



(d) Max Iterations: 3768

FIGURE 2.8: A magnified section of the Mandelbrot set image computed using an increasing number of maximum of iterations.

Chapter 3

Evaluation

This section gives an overview of ensuring the correctness of the output of the implementation.

3.1 Code Evaluation

In order to evaluate the correctness of the C code, I ran the main body of the code in modified evaluation program. This program functions by taking the screen co-ordinates from a memory address in the same way as if was the code was running on an FPGA, but rather than outputting the results to memory, the output is placed in a Portable Pixmap Format (PPM). PPM was chosen as it is an image file format ideal for use in Mandelbrot as images are stored as an array of pixel colour values. Mandelbrot correctness is best confirmed visually, if the code is even slightly wrong, the output will be far from similar to the correct image.

3.2 Simulation

The purpose of simulation is to verify the functionality and timing of a design by interpreting VHDL code as if it was circuit functionality. The logical results are then displayed on screen [26]. For the simulation I used XILINX's ISE Simulator, *ISIM*. Simulation makes use of a Test Bench, which is HDL code that provides a set of stimuli to the simulation[27].

To start, compiled and formatted instructions are inserted into the MIPS Instruction BRAM. This is seen in Figure 3.5, the comments beside the instructions show their meaning in assembly.

```

16 architecture a_instr_rom of instr_rom_BRAM_com is
17
18 type programm_sequence is array (integer range <>) of std_logic_vector(31 downto 0);
19 signal program : programm_sequence(0 to (2 ** size)-1) := (
20 X"27bdfc8", -- addiu sp,sp,-56
21 X"afbf0034", -- sw ra,52(sp)
22 X"afbe0030", -- sw s8,48(sp)
23 X"03a0f021", -- move s8,sp
24 X"3c1c0000", -- lui gp,0x0
25 X"279c0000", -- addiu gp,gp,0
26 X"afbc0010", -- sw gp,16(sp)
27 X"afc00028", -- sw zero,40(s8)
28 X"24020004", -- li v0,4
29 X"afc20024", -- sw v0,36(s8)
30 X"24020008", -- li v0,8
31 X"afc20020", -- sw v0,32(s8)
32 X"3c020000", -- lui v0,0x0
33 X"c4420004", -- lwcl $f2,4(v0)
34 X"00200825", -- move at,at
35 X"c4430000", -- lwcl $f3,0(v0)
36 X"3c020000", -- lui v0,0x0
37 X"c4400004", -- lwcl $f0,4(v0)
38 X"00200825", -- move at,at
39 X"c4410000", -- lwcl $f1,0(v0)
40 X"00200825", -- move at,at
41 X"46201081", -- sub.d $f2,$f2,$f0
42 X"3c020000", -- lui v0,0x0

```

FIGURE 3.1: A screen shot of the some of the cross-compiled assembly code in the the Instruction BRAM of the MIPS overlay.

After this ISIM is launched. In Figure 3.6 we see the button in ISE whcih launches ISIM. Before ISIM is launched, ISE performs a check of the behavioural syntax of the source file [28].

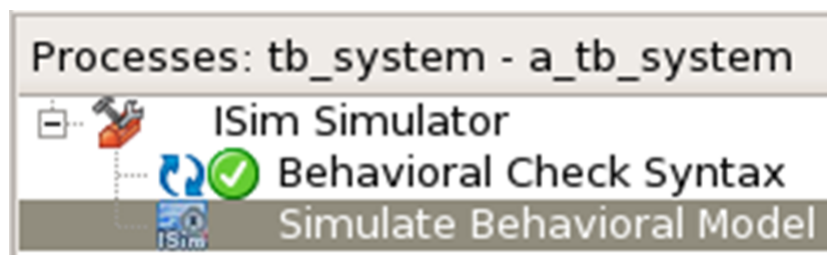


FIGURE 3.2: A screen shot showing the button which launches the simulator after checking the Behavioural Syntax.

Figure 3.7 shows the behavioural simulation running. In this image the scale is per microsecond. Verification is possible by taking note of the most important signals as seen in Figure 3.8. By manually checking what I expect the address of the next instruction to be, I can ensure that the design is functionally working and correct.

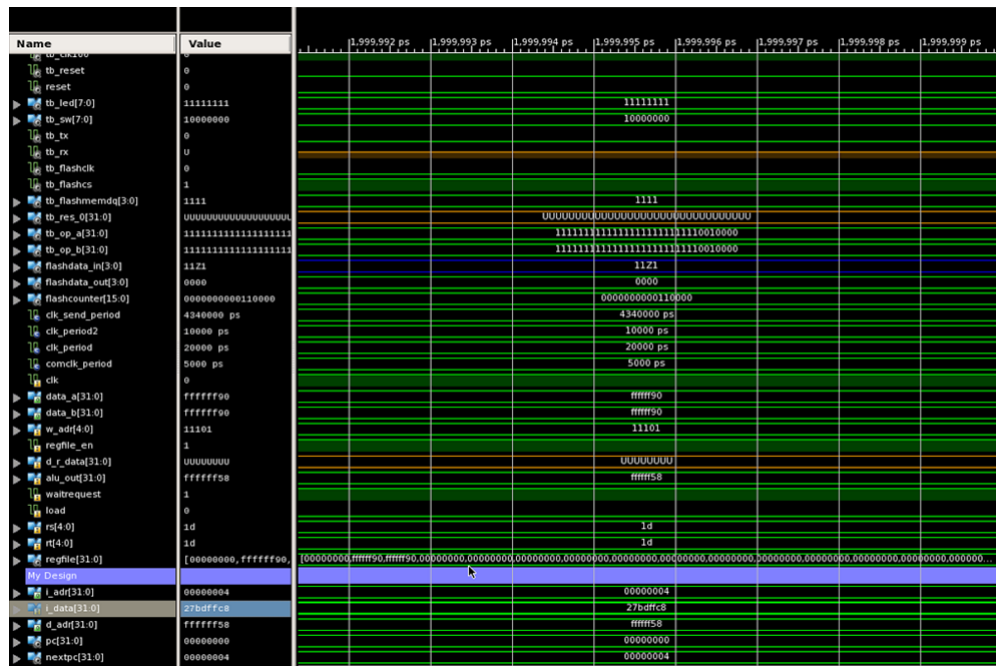


FIGURE 3.3: A screen shot of the simulation window.




FIGURE 3.4: A close up screen shot of the most relevant signals in the simulation.

Chapter 4

Reflections and Conclusion

4.1 Change of Objective

Initially, this project's objective was to render the Mandelbrot set on a dedicated FPGA Processing Element, rather than use a MIPS hardware overlay. After a few weeks of preliminary research, and after a discussion with my project supervisor, it was decided that the current project objective was a more worthwhile third year project. The modified project would be more applicable to hardware design in the real world, it would allow me to learn about a wider variety of topic areas and it would be flexible enough to be expanded upon (See chapter 5).

4.2 Challenges

Over the course of this project I faced a few roadblocks to my progress. While I had budgeted time for unexpected events, overcoming the hurdles took longer than expected. This had the unfortunate effect of slowing down all progress, as I could not continue with the later stages.

On university machines, undergraduate students are not able to gain superuser access for administrator privileges. As I was unable to gain admin privileges, I was unable to correctly compile the cross compiler toolchain. This was a crucial step as it allows both parts of the system (C code and MIPS overlay) to work together.

This was overcome by the use of an Oracle VirtualBox Virtual Machine [29] running from a 32GB Flash Drive. On this virtual machine I then installed a 64-bit Ubuntu version 14.04.4 (Long Term Support version). This granted me full admin privileges and I did all of my cross compilation on this virtual machine.

4.3 Conclusion

I was motivated to choose the project as I had held an interest in fractal images and their appearance in nature for a long time. I also saw the project as an opportunity to learn more about the topic area of hardware, as coming into the project, I had almost no computer hardware programming knowledge, let alone FPGA knowledge.

As could be expected, the learning curve for this project was relatively steep for me as hardware design differs significantly from software design. Furthermore, as FPGAs are not widely used consumer products, the information and resources available on this topic are rarely written for the audience of a newcomer to FPGA programming. Nonetheless I took this as a challenge and taught myself the basics of both popular hardware description languages, VHDL and Verilog to understand the existing code base.

Overall, I would describe this project as beneficial and worthwhile. Although I did not implement as much as I would have ideally liked to, I have enjoyed the process and I have learnt a great amount through research. I have tried to express this in this report.

Along with the technical knowledge gained in the implementation of a large project, I have also learnt equally valuable soft-skills such as organisation, time boxing and most importantly how to enthuse others about your ideas.

Chapter 5

Further Research

5.1 The Mandelbulb

First created in 2009 by Daniel White and Paul Nylander [30], the Mandelbulb is a 3-Dimensional hypercomplex fractal computed using spherical coordinates. In a similar fashion to the 2-Dimensional Mandelbrot image, the Mandelbulb can be zoomed in upon, revealing new unseen detail.



FIGURE 5.1: A render of the Mandelbulb [31].

5.1.1 Changes to be Made

In order to render the Mandelbulb on the current FPGA system, some changes would need to be made as the computation of the Mandelbulb differs from the Mandelbrot set in that it is 3-dimensional.

The Mandelbulb C code computes whether a given point in 3D space belongs to the Mandelbulb. By doing this a surface of the Mandelbulb is constructed. However, This is not enough to display the image on a screen. To display the image, we would need to use a Ray-marching, a simplistic 3D rendering technique which allows one to create an image of a surface by only knowing the distance from the surface to any point. Ray-marching would require some additional code to implement, but most of the structure of the system has been implemented.

Appendix A

Run Away to Infinity Criterion

Here I will show that if some z_n is farther than 2 from the origin, then successive iterates will grow without bound. That is, they will run away to ∞ [32].

For a complex number $z_n = x_n + iy_n$, the absolute value is

$$|z_n| = \sqrt{(x_n^2 + y_n^2)},$$

the distance from z_n to the origin.

Recalling the sequence z_0, z_1, \dots is defined by $z_{n+1} = z_n^2 + c$, we show if some z_n satisfies $|z_n| > \max(2, |c|)$, then the sequence z_n, z_{n+1}, \dots runs away to ∞ . So suppose $|z_n| > \max(2, |c|)$.

Because $|z_n| > 2$, we can write

$$|z_n| = 2 + \epsilon,$$

for some $\epsilon > 0$.

Now

$$|z_n^2| = |z_n^2 + c - c| \leq |z_n^2 + c| + |c|$$

So

$$\begin{aligned} |z_n^2 + c| &\geq |z_n^2| - |c| = |z_n|^2 - |c| \\ &> |z_n|^2 - |z_n| \text{ (because } |z_n| > |c| \text{)} \\ &= (|z_n| - 1) \cdot |z_n| = (1 + \epsilon) \cdot |z_n| \end{aligned}$$

That is, $|z_{n+1}| > (1 + \epsilon) \cdot |z_n|$. Iterating, $|z_{n+k}| > (1 + \epsilon)^k \cdot |z_n|$.

To complete the proof that $|z_n| > 2$ implies the sequence runs away to infinity, observe that if $|c| > 2$, then

$$z_0 = 0$$

$$z_1 = c$$

$$and z_2 = c^2 + c = c \cdot (c + 1)$$

$$so |z^2| = |c| \cdot |c + 1| > |c| (noting |c + 1| > 1 because |c| > 2).$$

Bibliography

- [1] John Hutchinson. Fractals and Self-Similarity. *Indiana University Mathematics Journal*, 30(5):713–747, 1981.
- [2] Nigel Lesmoir-Gordon. *The Colours of Infinity: The Beauty, The Power and the Sense of Fractals*. Clear Books, 2004. ISBN 1904555055.
- [3] Michael Frame and Nial Neger. Julia sets - computational issues. URL http://users.math.yale.edu/public_html/People/frame/Fractals/. [Accessed : 20/03/2016].
- [4] Frank Hannig Dirk Koch, Daniel Ziener. *FPGAs For Software Programmers*. Springer, 2016.
- [5] Charles Eric LaForest. *High-Speed Soft-Processor Architecture for FPGA Overlays*. PhD thesis, University of Toronto, 2015.
- [6] Ian Kuon, Russell Tessier, and Jonathan Rose. Fpga architecture: Survey and challenges. *Foundations and Trends in Electronic Design Automation*, 2(2):46, 2008. ISSN 1551-3939. doi: 10.1561/10000000005. URL <http://dx.doi.org/10.1561/10000000005>.
- [7] Wim Vanderbauwhede. *High-performance computing using FPGAs*. Springer, New York, NY, 2013. ISBN 978-1-4614-1790-3.
- [8] D. Koch, F. Hannig, and D. Ziener. *FPGAs for Software Programmers*. Springer International Publishing, 2016. ISBN 9783319264066.
- [9] Field programmable gate array. URL <http://www.xilinx.com/training/fpga/fpga-field-programmable-gate-array.htm>. [Accessed : 5/03/2016].
- [10] J. Kepner. *Parallel MATLAB for Multicore and Multinode Computers*. SIAM e-books. Society for Industrial and Applied Mathematics (SIAM, 3600 Market Street, Floor 6, Philadelphia, PA 19104), 2009. ISBN 9780898718126.
- [11] Nexys4 ddr artix-7 fpga board. URL <http://www.xilinx.com/products/boards-and-kits/1-6olhw1.html>. [Accessed : 13/04/2016].

- [12] H. Tucker Jerry. Hardware description languages.
- [13] M.D. Ciletti. *Advanced Digital Design with the Verilog HDL*. Pearson Education, 2011. ISBN 9780133002546. URL <https://books.google.co.uk/books?id=QTArAAAAQBAJ>.
- [14] F. Rodriguez-Henriquez, N.A. Saqib, A.D. Pérez, and C.K. Koc. *Cryptographic Algorithms on Reconfigurable Hardware*. Signals and Communication Technology. Springer US, 2007. ISBN 9780387366821.
- [15] Graham Schelle, Jamison Collins, Ethan Schuchman, Perry Wang, Xiang Zou, Gautham China, Ralf Plate, Thorsten Mattner, Franz Olbrich, Per Hammarlund, Ronak Singhal, Jim Brayton, Sebastian Steibl, and Hong Wang. Intel nehalem processor core made fpga synthesizable. In *Proceedings of the 18th Annual ACM/SIGDA International Symposium on Field Programmable Gate Arrays, FPGA '10*, pages 3–12, New York, NY, USA, 2010. ACM. ISBN 978-1-60558-911-4.
- [16] Perry H. Wang, Jamison D. Collins, Christopher T. Weaver, Bliappa Kuttanna, Shahram Salamian, Gautham N. China, Ethan Schuchman, Oliver Schilling, Thorsten Doil, Sebastian Steibl, and Hong Wang. Intel®atom™processor core made fpga-synthesizable. In *Proceedings of the ACM/SIGDA International Symposium on Field Programmable Gate Arrays, FPGA '09*, pages 209–218, New York, NY, USA, 2009. ACM. ISBN 978-1-60558-410-2.
- [17] A. Koch, R. Krishnamurthy, J. McAllister, R. Woods, and T. El-Ghazawi. *Reconfigurable Computing: Architectures, Tools and Applications: 7th International Symposium, ARC 2011, Belfast, UK, March 23-25, 2011, Proceedings*. Lecture Notes in Computer Science. Springer Berlin Heidelberg, 2011. ISBN 9783642194757.
- [18] Anders Hauk Fritzell. A system for fast dynamic partial reconfiguration using goahead: Design and implementation. pages 28,37, 2013.
- [19] D. Sweetman. *See MIPS Run*. The Morgan Kaufmann Series in Computer Architecture and Design. Elsevier Science, 2010. ISBN 9780080525235.
- [20] J.L. Hennessy and D.A. Patterson. *Computer Architecture: A Quantitative Approach*. The Morgan Kaufmann Series in Computer Architecture and Design. Elsevier Science, 2006. ISBN 9780080475028.
- [21] David Goldberg. What every computer scientist should know about floating-point arithmetic. *ACM Comput. Surv.*, 23(1):5–48, March 1991. ISSN 0360-0300.
- [22] Ieee 754 single floating point format. URL https://commons.wikimedia.org/wiki/File:IEEE_754_Single_Floating_Point_Format.svg. [Accessed : 7/02/2016].

- [23] Fixed point arithmetic. URL <http://radio.feld.cvut.cz/matlab/toolbox/filterdesign/quant14a.html>. [Accessed : 6/02/2016].
- [24] Brain Hall. The mandelbrot set. URL <http://beej.us/blog/data/mandelbrot-set/>. [Accessed : 25/04/2016].
- [25] Why do i need a cross compiler ? URL http://wiki.osdev.org/Why_do_I_need_a_Cross_Compiler%3F. [Accessed : 28/04/2016].
- [26] Simulation overview. . URL http://www.xilinx.com/itp/xilinx10/isehelp/ise_c_simulation_overview.htm. [Accessed : 27/04/2016].
- [27] Test benches. . URL http://www.xilinx.com/itp/xilinx10/isehelp/ise_c_simulation_test_bench.htm. [Accessed : 27/04/2016].
- [28] Checking syntax for simulation. . URL http://www.xilinx.com/itp/xilinx10/isehelp/pp_p_process_check_syntax_simulation.htm.
- [29] Virtualbox manual. URL <https://www.virtualbox.org/manual/ch01.html>. [Accessed : 28/04/2016].
- [30] Paul Nylander. Hypercomplex fractals. URL <http://bugman123.com/Hypercomplex/>. [Accessed : 20/04/2016].
- [31] Ondrej Karlik. Power 8 mandelbulb fractal overview. URL https://en.wikipedia.org/wiki/File:Power_8_mandelbulb_fractal_overview.jpg. [Accessed : 27/04/2016].
- [32] Michael Frame. Run away to infinity criterion. URL http://users.math.yale.edu/public_html/People/frame/Fractals/. [Accessed : 5/03/2016].